

Models selecting prototypes with ProtoPNet

ADAM KOZŁOWSKI

24.01.2022



Presentation overview

- Prototypes generation in XAI
- ProtoPNet architecture and training
- Experiment

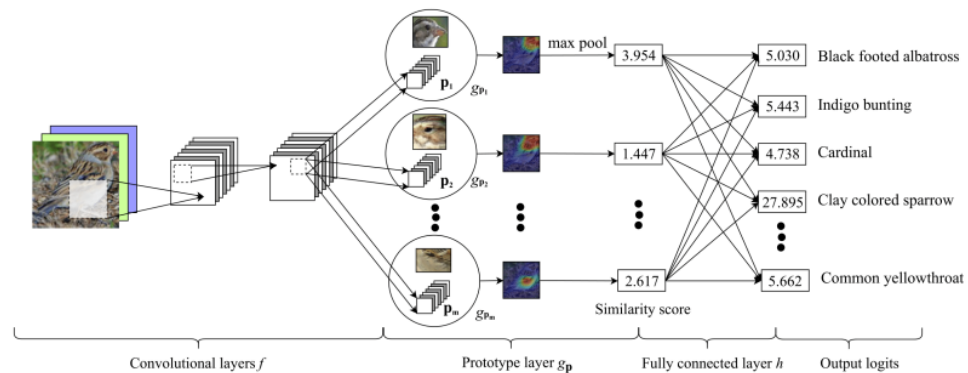




Figure 2: ProtoPNet architecture.


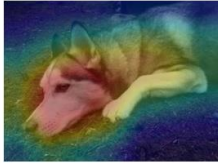



Prototypes in XAI

Previous XAI methods

Posthoc methods



- Used on already trained models
- The most important techniques are: activation maximalization, deconvolution, saliency visualization
- They don't actually explain models decision

Test Image	Evidence for Animal Being a Siberian Husky	Evidence for Animal Being a Transverse Flute
	 "Explanation"	

Źródło: IMIC Keynote 3, C. Rudin MICCAI 2021

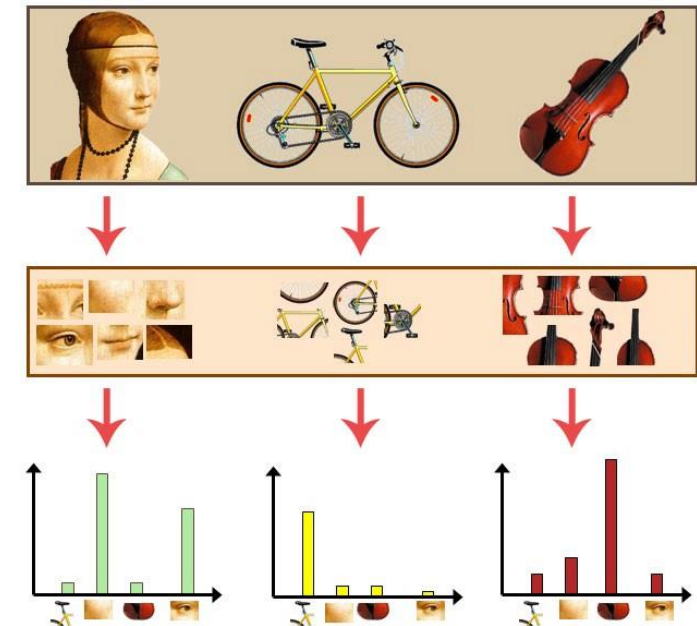

Methods with built-in attention mechanisms

- They mark areas that are important in models decision process
- The most prominent ones are: models with activation maps and part-based models
- They only mark the important area – they don't explain why



Prototypes

- Main inspiration is the way how human experts try to explain complex ideas with examples
- The model explains its decision by providing set of prototypes from training data
- Current methods are based on the concept of *bag-of-visual-words*
- So far prototype selection was separate from feature extraction



[Bag of Visual Words in a Nutshell | by Bethea Davida](#)
[| Towards Data Science](#)

ProtoPNet (*Prototype Part Network*)

- Architecture proposed in „This looks like that: Deep Learning for Interpretable Image Recognition” (Chen et al. 2019 [1806.10574.pdf \(arxiv.org\)](https://arxiv.org/pdf/1806.10574.pdf))
- Principles
 - Prototypes are chosen based on the distance to sample in latent space
 - Prototype features are used during training
 - Interpretable architecture
- Inspired by previous methods. The closest ones are: Bayesian Case Model (main idea) and method from „Deep Learning for Case-Based Reasoning through Prototypes” (Li et al. 2017 <https://arxiv.org/pdf/1710.04806>) which takes autoencoder for feature selection
- Architecture successfully used in IAIA-BL for detecting and evaluating breast lesions (Barnett et al. 2021, [2103.12308.pdf \(arxiv.org\)](https://arxiv.org/pdf/2103.12308.pdf))

ProtoPNet architecture and training

ARCHITECTURE, LOSS FUNCTION, PROTOTYPES SELECTION

Architecture

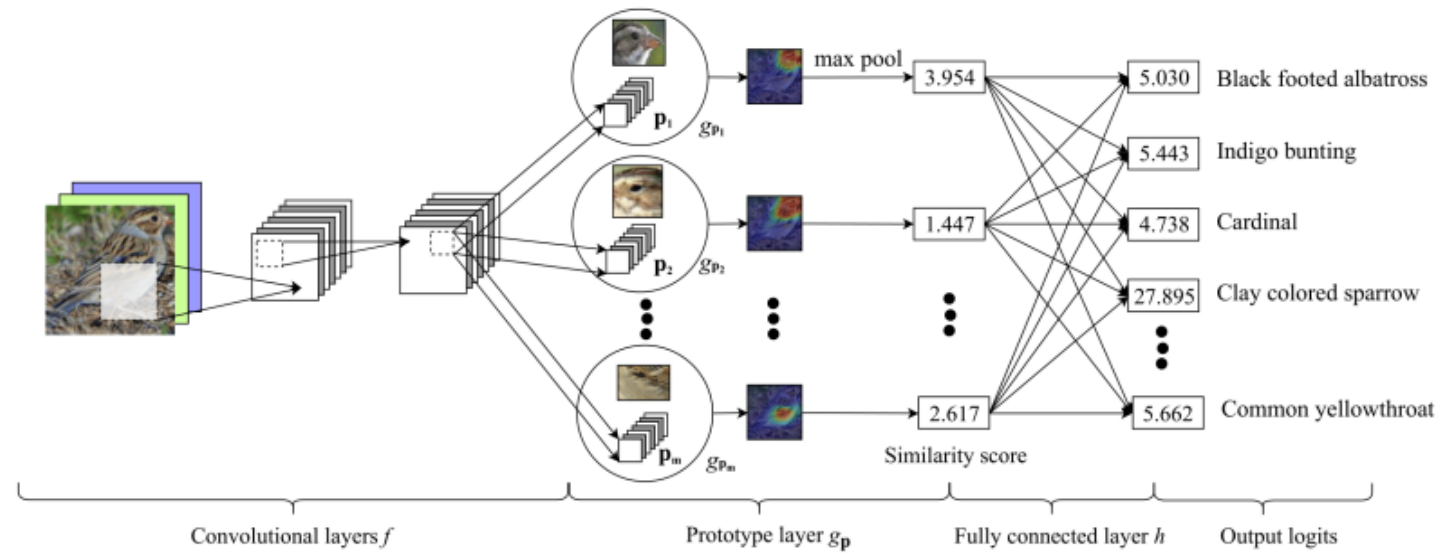
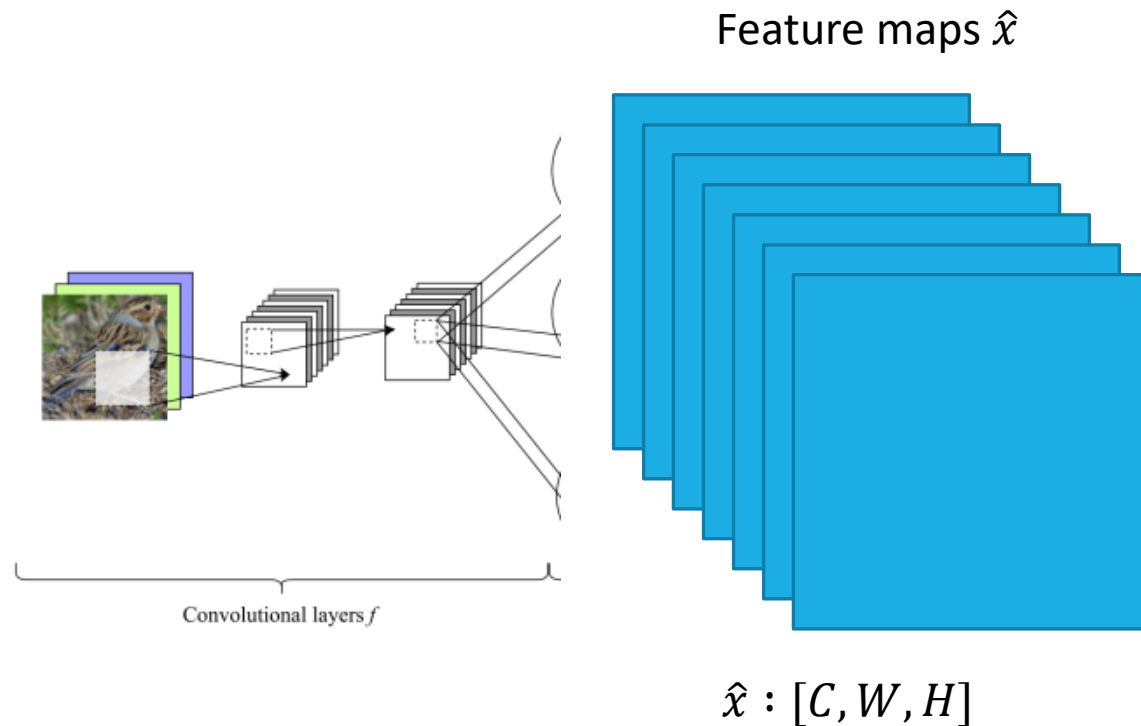


Figure 2: ProtoPNet architecture.

Latent feature maps




- For given input x convolutional layers f extract useful features
- Output $f(x)$ is passed through additional convolutional layers with 1×1 filters
- Number of channels of these layers corresponds to the size of prototype vectors length
- The final output is tensor \hat{x} of size

$$\hat{x} : [C, W, H]$$


C – num. of channels, W -width, H - height
example: using VGG16 the spatial dimension of convolutional output is $[512, 5, 5]$, then it is reduced with additional layers to e.g. $[C=128, 5, 5]$

Prototype part vectors

Prototype part vector p_j



$$p_j : [C, W_p, H_p]$$

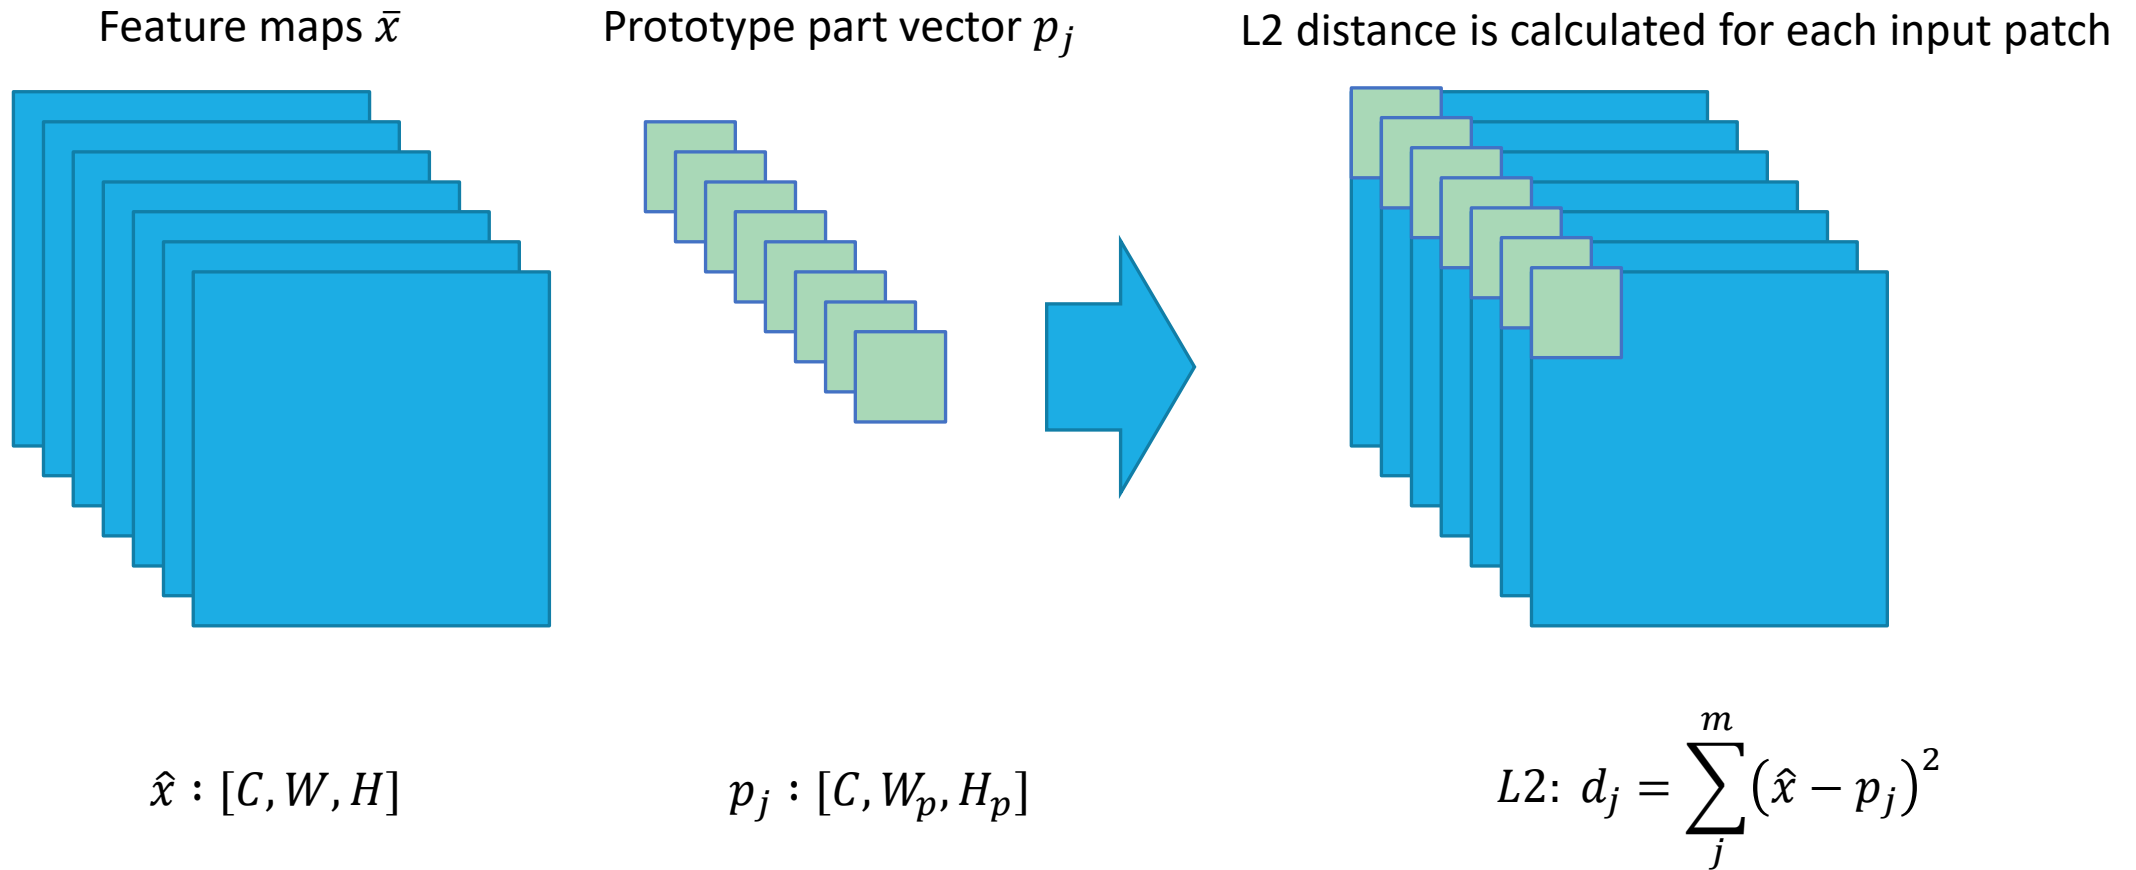


- Based on image prototypes selected during training
- They represent class specific features
- Each vector represents single part of the prototype image

$p_j : [C, W_p, H_p]$
 C – num. of latent channels, W_p – prototype width, H_p – prototype height
e.g. [512,1,1]

- Each of the dimensions is a hyperparameter

L2 distance to prototype part vectors




L2 distance to prototype part vectors

Feature maps \bar{x}




$$\hat{x} : [C, W, H]$$

Prototype part vector p_j




$$p_j : [C, W_p, H_p]$$



$$L2: d_j = \sum_j^m (\hat{x} - p_j)^2$$


L2 distance to prototype part vectors

Feature maps \bar{x}




$$\hat{x} : [C, W, H]$$

Prototype part vector p_j

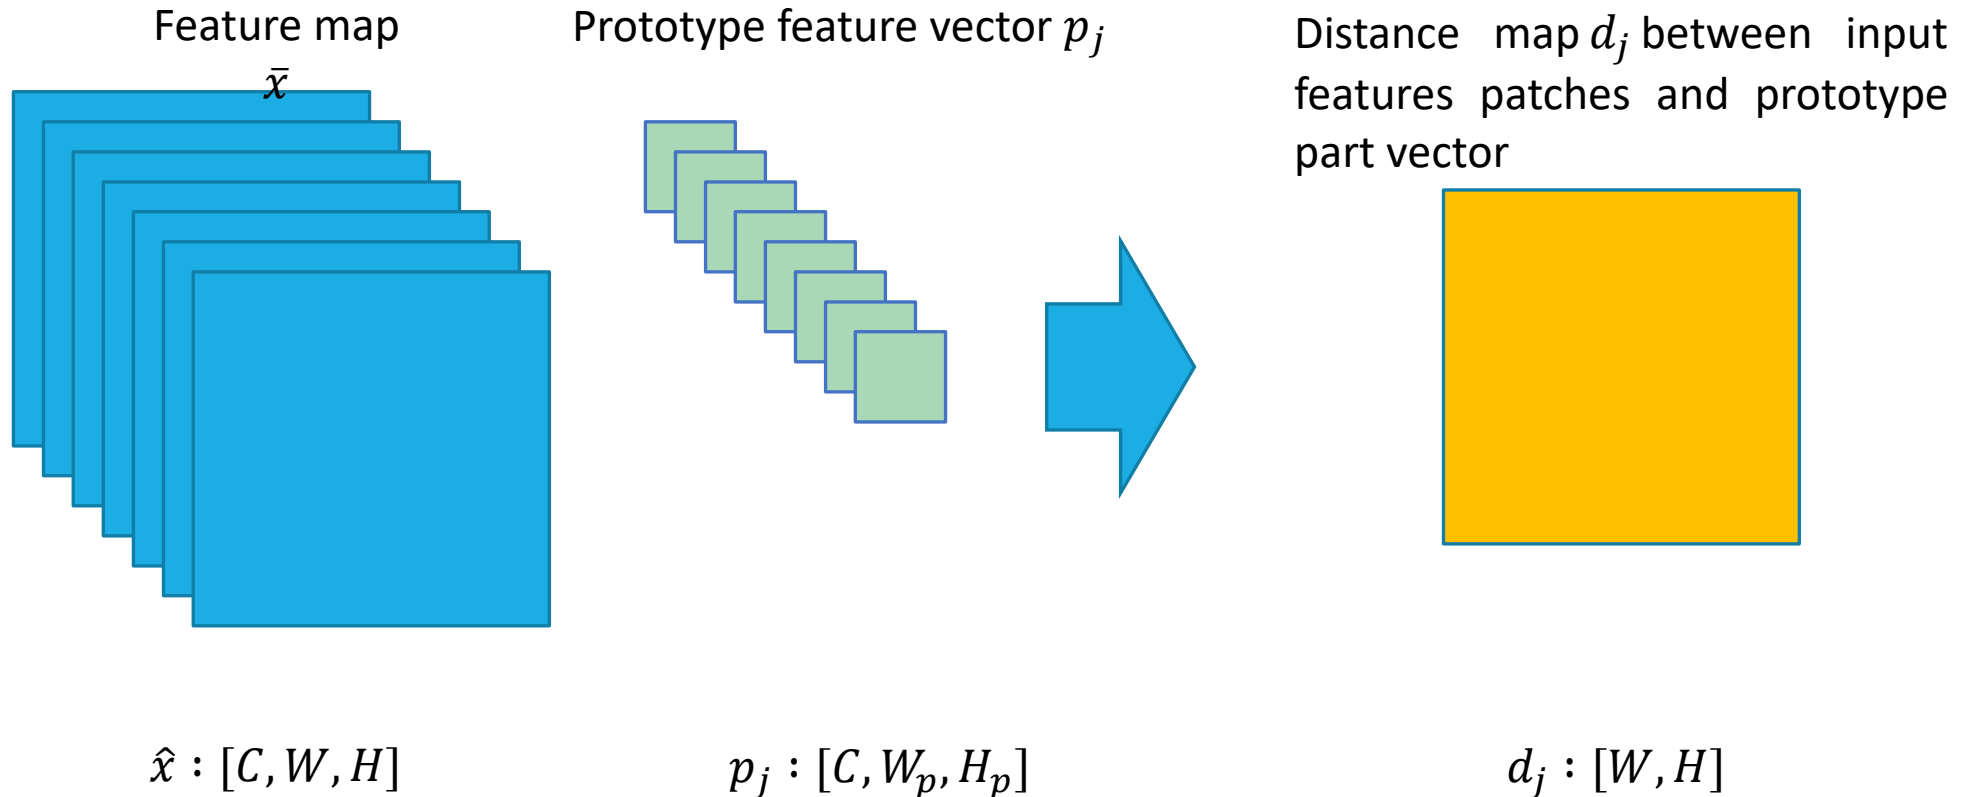


$$p_j : [C, W_p, H_p]$$




$$L2: d_j = \sum_j^m (\hat{x} - p_j)^2$$

L2 distance to prototype part vectors




Similarity score

Distance maps d to each
of the prototypes



Similarity score vector s



- Similarity score is calculated as:

$$s_j = \log \frac{\min d_j + 1}{\min d_j + \epsilon}$$

between input feature map and j -th prototype
part vector



Intuitively:

if $d_j \rightarrow 0$ then $s \rightarrow \log \frac{1}{\epsilon}$

when $d_j \rightarrow \infty$ then $s \rightarrow 0$

Classification

Similarity score s



- To get output predictions \hat{y} the similarity scores are passed through dense layer

Interpretability! – you can measure how similarity score to each prototype influenced prediction

- Dense layer weights initialized as:
 - 1 if output class is the same as prototype class
 - -0.5 otherwise

Training

The loss function for training is as follows:


$$loss = \min_{P, w_{conv}} \frac{1}{n} \sum_{i=1}^n CrsEnt(\hat{y}_i, y_i) + \lambda_1 Clst + \lambda_2 Sep$$

- $Clst$ is the shortest distance between sample feature map and prototype part vector of the same class
- Sep is the shortest distance to prototype part vector
- Authors suggest that $\lambda_2 = -0.1\lambda_1$ to enforce higher separation between classes

Convergence if:

1. Crossentropy is close enough to local minimum
2. $Clst < Sep$ which means closer distance between parts of prototypes belonging to the same class

Prototype selection




Prototype part vectors P


1. Prototypes are selected during training
2. Prototype part vectors are randomly initialized
3. Every few epochs prototypes are updated
4. Each class must have fixed number of prototypes

Prototype selction

For prototype p_j

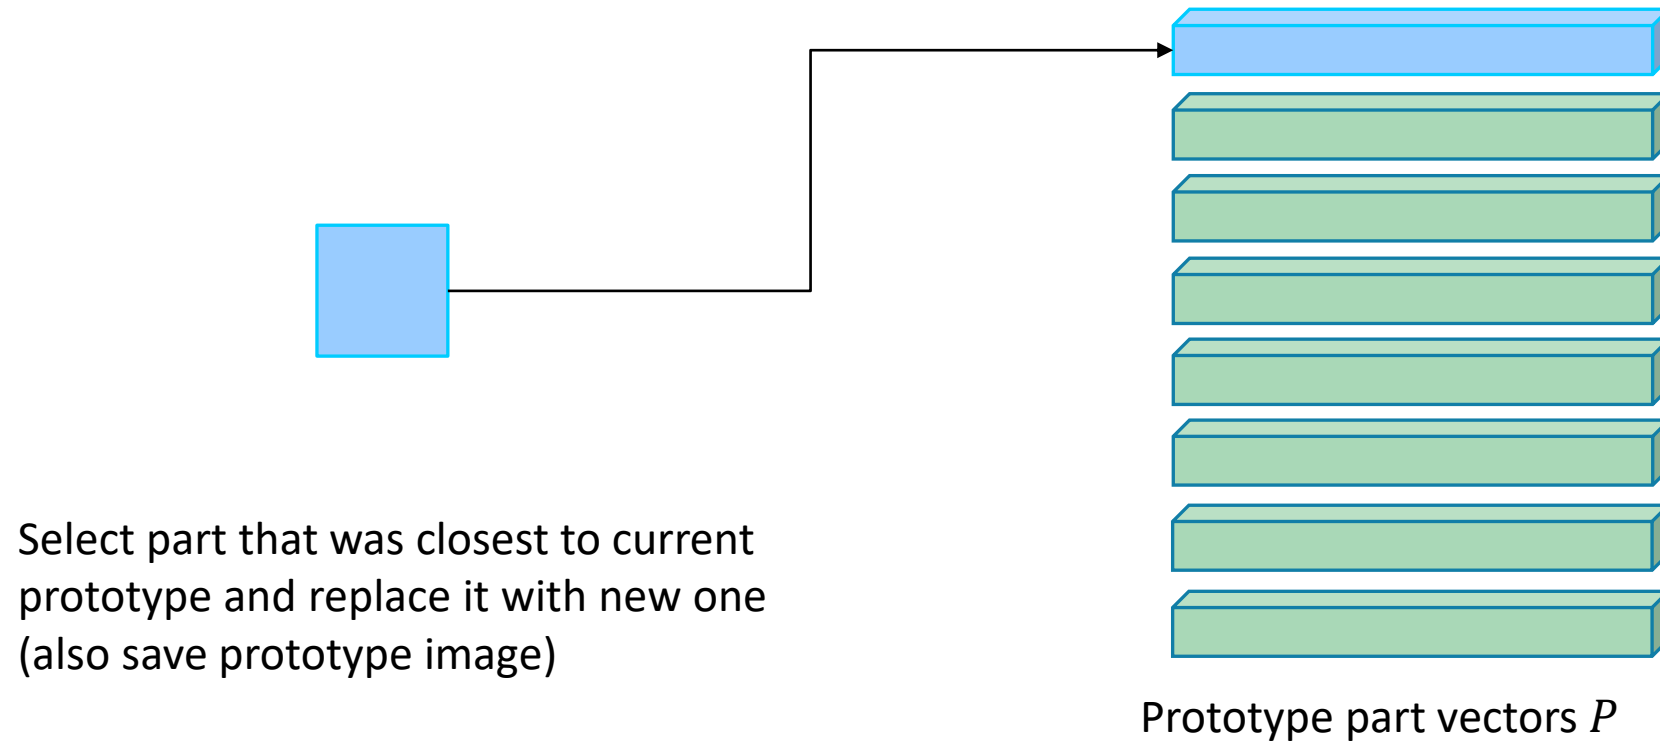


Calculate L2 distance to all the feature maps parts for the whole training data




Prototype part vectors P

Prototype selction



Prototype selction



Repeat the whole proces until all prtotype part vectors are updated

Prototypes part vectors P

Experiment



PROTOTYPES AND METRICS RESULTS

Experiment setup

Experiment goal was to recreate ProtoPNet model for different task and to explain it's decisions for given images


- Dataset: *Rock, paper, scissors* (balanced, 2892 images, augmented to 8676, 80-20 train and test split)
- Pretrained VGG16 convolutional layers for feature extraction
- 45 prototype part vectors (15 for each class) with dimensions [128,1,1]
- Metrics: accuracy and elements of the loss function (e.g. *Clst, Sep*)
- Trained for 100 epochs

Results - prototypes




class	top 3 prototypes						top incorrect prototype	
	similarity	connection	similarity	connection	similarity	connection	class 1	class 2
rock	0,5744	1,119	0,2074	1,0695	0,1937	1,0741	scissors; 0,1545	rock; 0,03635
paper	3,9593	1,0205	0,3525	1,0196	0,3503	1,0172	rock; 0,03622	scissors; 0,03355
scissors	2,0435	1,0909	0,2241	1,0839	0,2166	1,0869	rock; 0,01371	paper; 0,0137


Results - prototypes




Results - prototypes




Results - prototypes




Results - metrics



Results - metrics



Results - metrics



ProtoPNet - summary

Pros:

- Interpretable model output
- Informative visual explanation
- Universal – can be used in various different tasks
 - Unsupervised learning – prototypes can be non-class specific
 - Different input data – prototypes are selected based on latent features
 - Can be used with every architecture

Cons:

- Requires additional training
- Longer training time (higher computational cost + additional metrics to optimize)
- Additional hyperparameters (e.g. number of prototypes)

Thank you!
