

Hyp-RL

Hyperparameter optimization with reinforcement learning

Bartłomiej Sobieski

Faculty of Mathematics and Information Science
Warsaw University of Technology

May 23, 2022

Overview

1. What is reinforcement learning?
2. Q-learning
3. Reinforcement learning and hyperparameter optimization
4. Our approach

What is reinforcement learning?

How do children learn to walk?



Intuitive understanding

An **agent** lives in the **environment** that, at a given moment, can be described by its **state**. By taking an **action** in this environment , the agent receives a **reward** and changes the state of the environment.

Intuitive understanding

- Environment
- Agent
- States
- Actions
- Rewards

The agent's goal is to maximize the cumulative rewards it receives over time.

Markov decision process

A Markov decision process is characterized by:

- **S** - set of states
- **A** - set of actions
- **R** - set of rewards

For timestep $t = 0, 1, 2, \dots$

$$S_t \in S \rightarrow A_t \in A \rightarrow (S_t, A_t) \xrightarrow{t+1} f(S_t, A_t) = R_{t+1}, S_{t+1} \in S$$

Return and discounted return

Return:

$$G_t = R_{t+1} + R_{t+2} + \dots + R_T$$

Discounted return:

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \text{ for } \gamma \in (0, 1)$$

The agent's goal is to maximize the expected discounted return of rewards.

How does an agent act?

An agent acts according to a **policy** π .

$$\pi(A_t|S_t) = \mathbb{P}(A_t|S_t)$$

Q-learning

State values

$$v_{\pi}(S_t = s) = E_{\pi}(G_t \mid S_t = s)$$

Q-value function

$$Q_{\pi}(S_t = s, A_t = a) = E_{\pi}(G_t \mid S_t = s, A_t = a) =$$
$$E_{\pi} \left(\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s, A_t = a \right)$$

Optimal policy

$$\pi \geq \pi' \Leftrightarrow \forall_{s \in S} v_{\pi}(s) \geq v_{\pi'}(s)$$

$$\pi \text{ is an optimal policy} \Leftrightarrow \forall_{\pi'} \pi \geq \pi'$$

Bellman optimality equation


For an optimal policy π^*

$$Q_{\pi^*}(s, a) = \max_{\pi} Q_{\pi}(s, a).$$

Bellman optimality equation:

$$Q_{\pi^*}(s, a) = E \left(R_{t+1} + \gamma \max_{a'} Q_{\pi^*}(s', a') \right).$$

Deep Q-learning



Policy network

Deep Q-learning

$$Loss = Q_{\pi^*}(s, a) - Q(s, a)$$

$$\Leftrightarrow$$

$$Loss = E \left(R_{t+1} + \gamma \max_{a'} Q_{\pi^*}(s', a') \right) - E \left(\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \right)$$

Target network

$$Loss = E \left(R_{t+1} + \gamma \max_{a'} Q_{\pi^*}(s', a') \right) - E \left(\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \right)$$


↓

$$\max_{a'} Q_{\pi^*}(s', a')$$

↑

Target network

Target network



Target network

Collecting experience

$$e_t = (s_t, a_t, r_{t+1}, s_{t+1})$$

Replay buffer = $\{e_t : t \in T - \text{capacity}, \dots, T\}$,

where T is the current timestep number.

Exploration vs. exploitation

$$U \sim \mathcal{U}([0, 1])$$

$$\epsilon = \epsilon_{end} + (\epsilon_{start} - \epsilon_{end}) \cdot e^{-t \cdot \epsilon_{decay}}$$

$\epsilon < U \rightarrow$ action chosen randomly.

$\epsilon > U \rightarrow$ action chosen by the agent.

Reinforcement learning and hyperparameter optimization

Problem overview

$$M_{\lambda}(D_{train})$$

$$\lambda \in \Lambda = \Lambda_1 \times \dots \times \Lambda_p$$

$$\lambda^* = \arg \min_{\lambda \in \Lambda} \mathcal{L}(M_{\lambda}(D_{train}), D_{valid})$$

Find the best hyperparameter configuration for a given dataset.

Hyp-RL

$$\mathcal{D} = \{D^1, \dots, D^m\}$$

$$\mathbf{A} = \Lambda$$

$$\mathbf{R} = \{f(M_\lambda(D_{valid}^i)) : \lambda \in \Lambda, i = 1, \dots, m\}$$

$$\mathbf{S} \ni s = (s_{static}, s_{dynamic})$$

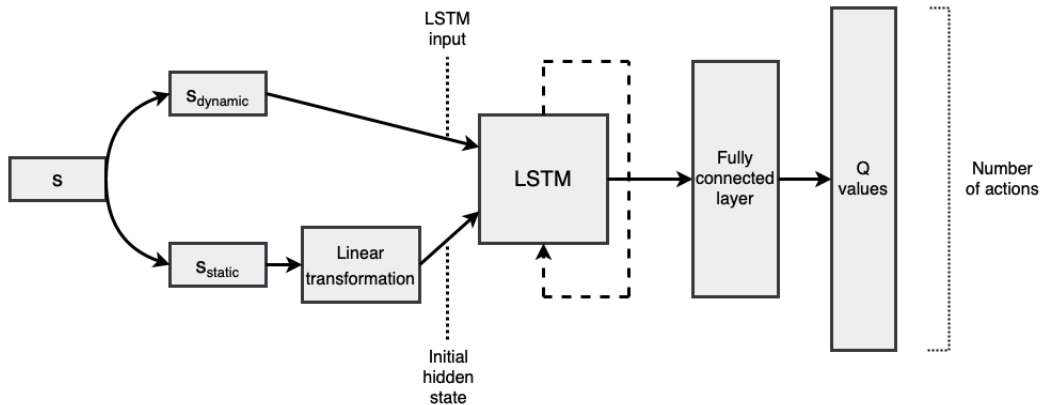
State decomposition

$$s = (s_{static}, s_{dynamic})$$

$$s_{static} \in \{\text{metadata}(D^i) : i = 1, \dots, m\}$$

$$s_{dynamic} \in (\Lambda \times R)^t \text{ for } t = 1, \dots, T$$

Policy network architecture



Algorithm

Algorithm Hyp-RL

1: **Input:** \mathcal{D} - set of datasets, Λ - hyperparameter grid, γ - discount factor, N_{target} - target update frequency, N_{replay} - replay buffer capacity, N_e - number of episodes per dataset, T - number of actions per episode.

2: Initialize policy network Q_{policy} parameters randomly and make Q_{target} as its clone, create replay buffer $\mathcal{B} = \emptyset$.

3: **for** $N_e \cdot |\mathcal{D}|$ **do**

4: Choose dataset D^i randomly

5: $s_t = (s_{static}(D^i), s_{dynamic} = (\{0\}^{P+1}))$

6: **for** $t = 0, \dots, T$ and while s_t is not terminal **do**

7: Determine next action as $a_t = \begin{cases} \sim \mathcal{U}(\Lambda) & p \sim \mathcal{U}([0, 1]) < \epsilon \\ \arg \max_a Q_{policy}(s_t, a) & \text{otherwise} \end{cases}$

8: Receive reward $r_t = f(M_{\lambda=a_t}(D_{valid}^i))$

9: Generate new state $s_{t+1} = s_t \cup \{(\lambda = a_t, r_t)\}$

10: Store new experience: $\mathcal{B} = \mathcal{B} \cup \{(s_t, s_{t+1}, a_t, r_t)\}$, replace oldest element if $|\mathcal{B}| > N_{replay}$

11: Sample a batch B of experiences from the replay buffer \mathcal{B} and relabel it as

12: $B = \{(s, a, Q(s, s', a, r)) \mid (s, s', a, r) \sim \mathcal{U}(\mathcal{B})\}$, where

13: $Q(s, s', s, r) = \begin{cases} r & s' \text{ is terminal} \\ r + \gamma \max_{a'} Q_{target}(s', a') & \text{otherwise} \end{cases}$

14: Update Q_{policy} by minimizing


15:
$$\sum_{(s,a,a') \in B} (Q - Q_{policy}(s, a))^2$$

16: Replace Q_{target} parameters with Q_{policy} parameters every N_{target} steps

17: **end for**

18: **end for**

How to use it?





Our approach

Hyperparameter grid

Hyperparameter	Type	Lower	Upper	Distribution
n_estimators	integer	1	1000	U
learning_rate	float	0.031	1	2^U
booster	discrete	-	-	U
subsample	float	0.5	1	U
max_depth	integer	6	15	U
min_child_weight	float	1	8	2^U
colsample_bytree	float	0.2	1	U
colsample_bylevel	float	0.2	1	U

1000 independant random configurations

metaMIMIC data



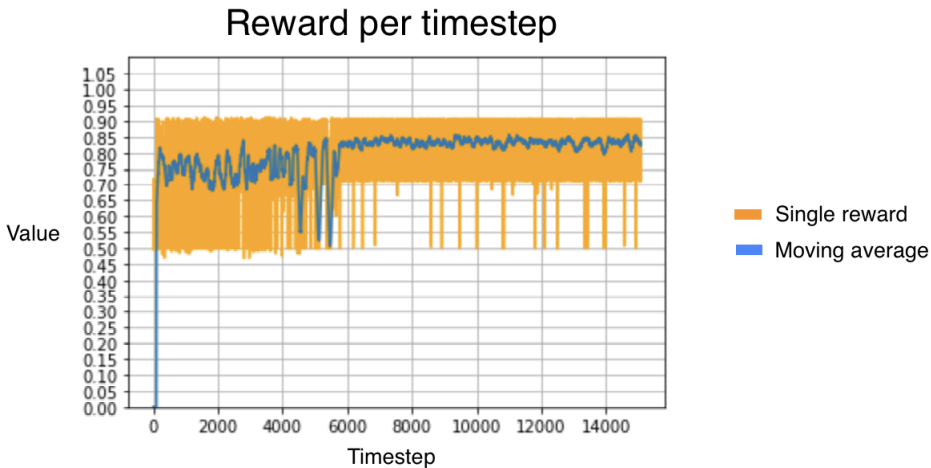
Value of ROC AUC for each hyperparameter configuration and target

Details

Reward \rightarrow ROC AUC

s_{static} \rightarrow Target ID

How should it look like?



Thank you for your attention